

The curled torus burps: activated escape and conversion in a graph model of emergent geometry

Emily Smith^{1,*}

¹*Independent researcher, Charlottesville, Virginia, USA*

(Dated: 10 October 2026)

We study how a curled torus gives way in four-regular bipartite graph energies next to combinatorial quantum gravity (CQG). A coefficient λ multiplies the price of a surplus square; only $\lambda = 1$ is CQG, and our decay experiments use $\lambda > 1$, where the ideal curled torus lies exactly $4(\lambda - 1)$ per vertex above the zero-energy states. For a specified Metropolis edge-switch chain, counting the exits offered per sweep predicts the mean first-exit time with no fitted parameter, consistently with direct measurements. During conversion most vertices keep a curled-like or flat-like square count, and one connected flat-like region usually dominates at half conversion, though several early patches occur. All 120 saved endpoints at $\lambda = 1.05$ are certified tori; growth of both periods is not established. At $\lambda = 1.25$, energy-conserving runs with a reservoir escape on a 12-unit budget over $N = 48$ –192, and exhaustive neutral-class enumeration certifies 12 as the lower bound at those sizes, while the available release grows with N . What the release produces depends on the reservoir’s thermal response; full release requires a zero-energy endpoint. All three registered conversion maps remain inconclusive. These results support a bounded kinetic claim and leave the conversion mechanism and reservoir equilibration open.

I. INTRODUCTION

Combinatorial quantum gravity (CQG) [1–3] treats space as a network: a discrete Ricci curvature, summed over a graph, favors configurations with geometric structure. Published equilibrium studies ask how geometry emerges from a random phase. Their ensembles require care: Ref. [2], Sec. 4, describes a restricted bipartite simulation in which no edge has surplus squares, whereas Ref. [3] discusses the uncapped curvature action, and Ref. [4] finds first-order behavior in a related global-only model without the same fixed bipartite ensemble. These are context, not identical-model validation. Kelly’s thesis [5], p. 219 and Fig. 4.58, reports small four-regular hysteresis loops over $N = 50$ –500 but declines a definitive first-order inference.

We ask instead how a specified ordered precursor escapes and converts at fixed coupling. The precursor is a tube: a torus curled to a period of four edges in one direction. It is stable for now in the energetic sense: at the sizes enumerated (Sec. II) no single move out of the perfect torus lowers its energy for the λ studied, yet lower states exist. We call its energy-releasing conversion a “burp.” “Decompactification” here means only removal of the four-cycle signature of the curled direction; growth of both geometric periods has not been measured (a $6 \times L$ torus already has the flat local signature with a fixed short period). This is a deformation next to CQG, not a CQG allotrope in the sense of Ref. [6].

Why deform CQG at all? At $\lambda = 1$ an edge carrying three squares costs exactly what an edge carrying two costs, so the flat torus, the curled torus and the 4-cube are degenerate: conversion between these ideal states re-

leases no energy (equal energies need not imply equal free energies). The coefficient λ lifts that degeneracy and changes nothing else: the states, the moves and the hard-core rule are those of CQG. For $\lambda > 1$ the ground states are the arrangements with two squares on every edge, a condition that fixes neither connectedness nor topology, and each curled direction costs a known energy per vertex. The family is therefore a controlled laboratory for a change between two ordered arrangements in which the barrier, the release and the relic can each be counted exactly, and $\lambda \rightarrow 1$ connects it to CQG.

Every kinetic statement below (waiting times, barriers, fronts) is about one Markov chain: the local edge switch used to sample the model [2], with Metropolis acceptance and an ordered-pair proposal that are our choices (Ref. [3] uses Glauber acceptance). Detailed balance gives the Boltzmann stationary measure but not mixing at production sizes, and we do not claim that the chain is the physical time evolution of emergent space. The energetics do not depend on it: the gap between the two arrangements, the cost of the cheapest single move out, and hence the least energy a sealed system needs to leave by such moves. The rates do, including how the waiting time scales with size (Sec. II).

A plain-language edition with diagrams accompanies the source (Methods); its chapters are cited as “PL, Ch. n ” where a concept is developed at greater length.

II. MODEL AND EXACT RESULTS

Diagrams for this section are in PL, Chs. 2 and 3.

A state is a simple undirected four-regular graph on a fixed labelled bipartition, $N/2$ vertices on each side, obeying the hard-core rule: no two vertices share more than two neighbours [2, 3]. Connectedness is not enforced. The energy is a price list on squares (four-cycles).

* emilysmithcreate@gmail.com

With S the number of squares, S_e the number on edge e , and $X = \sum_e (S_e - 2)_+$ the surplus squares on edges carrying more than two,

$$H = 16(N - S) + 4\lambda X = 4 \sum_e [(2 - S_e) + \lambda(S_e - 2)_+], \quad (1)$$

so a missing square costs 16 and a surplus square 4λ . The hard-core rule limits $S_e \leq 3$. At fixed dimensionless coupling $g > 0$, which acts as a temperature in the units of H , the target is $\pi(G) \propto e^{-H(G)/g}$, and a valid symmetric proposal is accepted with probability $\min(1, e^{-\Delta H/g})$. At $\lambda = 1$ this is the CQG curvature action: the convention $H = -4 \sum_v \sum_{w \sim v} \kappa(v, w)$ of Ref. [3], Eq. (21), with $\kappa(v, w) = -(2 - S_{vw})_+/2$ on these bipartite graphs, gives $H = 4 \sum_e (2 - S_e)_+$, each undirected edge counted once. For $\lambda \neq 1$, H is not this curvature action.

Zero is the floor (exact). Rewriting Eq. (1) as

$$H = 4 \sum_e [(2 - S_e)_+ + (\lambda - 1)(S_e - 2)_+]$$

shows $H \geq 0$ for $\lambda \geq 1$: no state lies below flat. For $\lambda > 1$, $H = 0$ if and only if every edge carries exactly two squares; at $\lambda = 1$, at least two. Flat lattice tori are ground states, and so are disconnected unions of them, so topology must be checked separately.

The ladder (exact). Periodic bipartite tori need even periods. Curl one direction of a torus to period four and the walk around it becomes a square, which puts a third square on every edge of each wrap-around 4-cycle. For even $L > 4$ the $4 \times L$ torus therefore has $S = 5N/4$ and $X = N$, and sits $H/N = 4(\lambda - 1) \equiv \varepsilon$ above a zero-energy state. Curling the other direction too gives the 4-cube (4×4), with $H/N = 8(\lambda - 1)$: each curled direction costs $4(\lambda - 1)$ per vertex in these ideal products (Fig. 1). A decay that ends with residual graph energy H_f releases $Q = \varepsilon N - H_f$, which is εN only for a zero-energy endpoint.

Counting signatures. A vertex has six pairs of incident edges. Let $d(v)$ be the number of pairs that close no square: 2 on a flat lattice torus, 1 on a one-curved torus, 0 on the 4-cube. This is not a geometric dimension. Under the hard-core rule a pair closes at most one square, so with q_v squares through v , $d(v) = 6 - q_v$ and $\sum_v d(v) = 6N - 4S$. The recorded conversion coordinate is

$$f = \frac{5N/4 - S}{N/4} = \bar{d} - 1,$$

which equals the fraction of vertices at $d = 2$ only when no signature other than $d = 1$ and $d = 2$ is present. Melting raises f too, and even $d = 2$ does not certify flat order (a valid single switch at $N = 64$ produces vertices whose incident edges carry $[1, 2, 2, 3]$ squares, not $[2, 2, 2, 2]$). We therefore call $d = 1$ curled-like and $d = 2$ flat-like counting signatures.

Other stuck states at $N = 18$ (exact enumeration). Enumerating all 26 graph classes and solving the affine inequalities $H > 0$ and $\Delta H > 0$ for every valid single

switch, with no grid in λ , gives precisely two strict local minima above the ground state on $\lambda \geq 1$: $(S, X) = (21, 12)$ on the open interval $1 < \lambda < 8/5$ and $(22, 16)$ on $1 < \lambda < 4/3$. This small-system certificate is not a large- N phase diagram.

The way out: counting the exits. The chain's one move swaps one partner each between two vertices of the same side (an edge switch preserving the bipartition), accepted by the Metropolis rule with the hard-core rule enforced. A proposal draws two edges one after the other, each uniformly from the $2N$ edges (a vertex of one side and one of its four slots). There are $4N^2$ equally likely ordered proposals, and every switch is proposed in two ways. A sweep is $2N$ attempts.

From the perfect $4 \times L$ torus, at the sizes enumerated below, the cheapest switches are of two kinds. Move A ($\Delta S = -2$, $\Delta X = -4$) costs $32 - 16\lambda$ and can be made in $3N$ ways, hence $6N$ ordered proposals; move B ($\Delta S = -4$, $\Delta X = -10$) costs $64 - 40\lambda$ and can be made in $2N$ ways, $4N$ proposals. A sweep therefore offers move A $2N \times 6N/4N^2 = 3$ times and move B twice, whatever N . Weighting each offer by its acceptance gives a two-channel approximation to the mean first-exit time with nothing fitted,

$$\tau_{AB}(\lambda, g) = [3e^{-(32-16\lambda)/g} + 2e^{-(64-40\lambda)/g}]^{-1}, \quad (2)$$

with A the cheaper for $\lambda < 4/3$ and both costs positive in the kinetic range used here. For a constant full per-attempt hazard p the exact first-exit law is geometric, with mean $1/(2Np)$ sweeps; the exponential is its small- p approximation.

First, $N/2$ switches change neither S nor X : each reconnects two neighbouring columns at opposite points, twisting the torus globally without leaving its energy. Breadth-first enumeration exhausts the neutral classes reached this way at $\lambda = 1.25$ for $N = 48, 64, 96, 144, 192, 288$. At each size there are three side-preserving isomorphism classes, all with the same full move census and minimum exit cost 12. Every valid switch is checked at each representative and every neutral target belongs to the exhausted list; since relabellings map switches bijectively, this certifies closure at these sizes, not for arbitrary L . Second, Eq. (2) leaves out every dearer exit. Each costs at least 10 more than the cheaper of A and B for $1.05 \leq \lambda \leq 1.45$, but those that break two distant edges at once are offered a number of times per sweep that grows in proportion to N (87 at $N = 192$). Counted exactly, they add at most 2% to the exit rate at the $g = 1.5$ conditions of the decay maps, and 16% at the hottest point of Fig. 2(a) ($g = 2.5$, $N = 144$), within its error bar. Third, that Eq. (2) does not depend on N is a property of the proposal, not of the barrier: the chain draws its two edges from the whole graph, so any one region is offered a way out about $1/N$ as often per sweep. A dynamics that updated every region at its own fixed rate would give a waiting time falling as $1/N$.




FIG. 1. (a) The flat lattice torus. (b) The $4 \times L$ torus drawn as a tube: its short direction is a single square, so the edges round each column carry three squares. (c) Schematic of the change of Sec. V, not a simulation snapshot: a possible boundary between curled and flat regions. Neither the schematic nor the square counts establish the number of moving interfaces. The graph ensemble also permits disconnected states.

TABLE I. Observables and clocks. Counts and components are measured on the evolving graph. The terms “first exit,” “detected decay” and “conversion” are not interchangeable.

Observable	Operational definition	Prediction or limitation
First exit T_1	First accepted move out of $S = 5N/4$, $X = N$; T22 counts each attempt. The exploratory Arrhenius run samples at sweep ends.	Eq. (2) is the A+B approximation; full constant hazard gives a geometric law in attempts.
Successful departure	Last departure before the commitment event $f \geq 0.25$; T22 counts exits and returns until that event or its 100,000-sweep cap.	239/240 reach commitment. Reciprocal mean exit count is a conditional empirical success statistic.
Detected decay W	First five-sweep check after 200 with $S/N < 1.25 - 3 \max(\sigma_{\text{rest}}, 10^{-6}) - 10^{-9}$; σ_{rest} uses the first forty samples.	Earliest value 205; threshold is trajectory-dependent. The floor curve alone does not reproduce this detector.
Conversion snapshot	First sampled $f \geq 0.25, 0.50, 0.75$ after the resting window, where $f = 5 - 4S/N$.	Histograms of d and induced $d = 2$ components refine the square count; neither certifies a geometric front.
Energy release	$Q = H_i - H_f$ for a saved endpoint; separately, the CSV release is H_i/N minus the final window mean of H/N .	Full ideal decrease requires $H_f = 0$; a window mean and a saved configuration need not agree.
Persistence	More than half a cell has $f_{200} < 0.25$.	A finite-duration criterion at $g = 1.5$, not an energetic spinodal.

III. THE COST OF LEAVING

At $\lambda = 1.25$ the cheapest move out of the perfect curled torus, move A, costs $\Delta E^\ddagger = 12$. This opens the door: it is a first-exit cost, not a proven minimum of the maximum energy along a complete conversion path, and whether an excursion then grows or falls back is a separate question (Sec. IV). The neutral-class enumeration of Sec. II makes 12 a lower bound at the enumerated sizes; an energy-conserving test asks whether it is also enough (PL, Ch. 4).

To seal the system we replace the bath at fixed g by a Creutz reservoir [7]: C stores of nonnegative energy. Each valid candidate move selects one store uniformly; the move is refused if that store cannot pay ΔH , and

otherwise the store’s energy changes by $-\Delta H$, so that $E_{\text{tot}} = H + \sum_i E_i$ is conserved. Stores and edge proposals are both selected globally: a seed is a total activation budget, not energy deposited at a chosen vertex. The recorded maximum drift of E_{tot} is zero in all T9 runs at $\lambda = 1.25$ (an observed check, not a claim about bitwise arithmetic at arbitrary λ).

A push that does not grow with size (exploratory). With one store initially holding $s \in \{4, 8, 10, 11, 11.5, 12, 13, 16\}$, the torus left its initial (S, X) values in 8 of 8 runs at each tested $s \geq 12$ and in 0 of 8 at each $s \leq 11.5$, for each $N = 48, 64, 96, 144, 192$ (320 runs of 12,000 sweeps). The lower half of this staircase is exact at these sizes: by the neutral-class certificate of Sec. II all nonneutral moves cost at least 12, so non-

negative stores with total budget below 12 cannot leave the basin. The upper half is an observation, dependent on protocol and duration. Eight successes give a two-sided exact 95% binomial interval of $[0.631, 1]$ at an individual setting (zero of eight gives $[0, 0.369]$), and the runner checks (S, X) once per sweep, so its success flag measures escape, not completed conversion. The push is fixed while the available ideal-state energy decrease εN grows with the system.

IV. THE FIRST EXIT

Table I defines the clocks used below, which are not interchangeable; PL, Ch. 5, develops the rate count.

Testing the count (exploratory). The prediction was written into the run's configuration before it ran, but not pre-registered. Over twelve conditions ($g = 1.4$ to 2.5 , $N = 64$ and 144 , 16 runs each) and a ninety-fold range of times, the ratio of the measured mean waiting time to the move-A term of Eq. (2) has a plain mean of 0.98, every ratio between 0.71 and 1.62 [Fig. 2(a)]. A weighted straight-line fit of $\ln \tau$ against $1/g$ gives effective slopes of 14.4 ± 1.0 at $N = 64$ and 12.1 ± 0.9 at $N = 144$; the first differs from 12 by 2.4 quoted standard errors. We do not read these slopes as independent measurements of a collective barrier: multiple channels give curvature in this plot, inverse-variance weighting based on 16 skewed waits can favor shorter sample means, and the scan reuses each replica's seed across couplings, so its fitted errors and pooled tests do not account for dependence between cells. The direct first-exit comparison below is the main rate test; model-conditional intervals are given in the Supplemental Material. Including move B lowers the predicted waiting time by 15% at $g = 1.5$; across λ , the full Eq. (2) is compared separately with first exits and detected-decay waits below.

The direct count (T22). Forty further runs in each of six cells ($\lambda = 1.05, 1.10, 1.25$; $N = 64, 96$) count exits attempt by attempt. All 240 first exits are observed, and their means are consistent with Eq. (2) in the model-conditional 95% intervals in the Supplemental Material. An opened door is not a departure, though: many exits fall back into the curled torus. Of the 240 runs, 239 reach $f \geq 0.25$ before the cap, and for these completed runs the reciprocal of mean exits per run is 0.56–0.68, with replica-bootstrap intervals spanning roughly 0.45–0.78 across cells. This estimates successful departures per counted exit under the $f \geq 0.25$ commitment rule. It is not a derived transmission coefficient: multiplying the exit rate by this fraction would require independent renewal excursions and control of time spent in failed excursions and growth.

A memoryless wait? (exploratory). A constant chance of exit per sweep gives a memoryless, exponential wait. The coefficient of variation is a coarse diagnostic, not sufficient to establish memorylessness. Dividing each time by its condition's mean and pooling over conditions, the

432 first exits we recorded give a Kolmogorov–Smirnov distance $D = 0.035$ from the unit exponential (nominal $p = 0.53$; shared-stream caveat in Methods) [Fig. 2(c)]. The 240 of them counted move by move (T22, above) are consistent with an exponential using the mean of Eq. (2), nothing estimated ($p = 0.85$). This does not establish exponential tails in each cell. Detected-decay waits show stronger departures, including the later T38 tail-test failures (Sec. VII).

V. THE CONVERSION

The decay itself was pre-registered: at $\lambda = 1.25$, $g = 1.5$, thirty decays at each of $N = 64, 96, 144, 192$, with the adjacency read when the conversion fraction first crossed 25, 50 and 75%. The predictions, averaged across replicas, were a memoryless waiting time (coefficient of variation between 0.7 and 1.3) and, at half conversion, at least 80% of vertices at $d \in \{1, 2\}$ and the converted vertices in one piece holding at least 70% of them. All three component tests hold at every size and on two independent sets of seeds (Table II). They show a dominant connected flat-like region in the evolving graph at half conversion, not unique nucleation or a single moving interface. The measured means show no clear decrease with N over these sizes; detected waits need not have the same size dependence as first exits.

The same pre-registration listed $\lambda = 1.5$; there the torus does not persist under the tested protocol at $g = 1.5$ (its number of squares per vertex falls to 0.95, below the flat torus's 1, within 50 sweeps in every replica), and no verdict was applied. A separate pre-registered test of whether a coarse law of a single front governs these decays (T12; 64 decays at $g = 1.5$ and 2.0 , under Metropolis and Glauber acceptance) came out NOT ESTABLISHED: 21 of the 59 that nucleated converted in more than one patch, mostly at the warmer coupling, and the front's rate fell roughly as $N^{-0.7}$.

The 25 and 75% snapshots were recorded but are not part of the registered criteria. They give the same picture: the mean share at $d \in \{1, 2\}$ is 0.988–0.997 and 0.985–0.995 across the eight cells, against 0.988–0.998 at 50% (PL, Ch. 6).

What is released. The energy difference between the ideal states is exact; what a given decay releases depends on where it ends. After settling, 24 to 28 of 30 decays per size have a window-averaged release within 1% of ε per vertex. This energy criterion alone does not certify torus topology. The other endpoints hold defects, read from their wiring: most often combinations of small defects of 8 to 45 units, and less often the four-point relic of Sec. VI, a closed loop of four vertices at $d = 1$ with $S = N + 1$ and, in its commonest form, $X = 6$, costing $24\lambda - 16$ (14 at $\lambda = 1.25$). Whether these states are stable or only slow is not known. The energy criteria were written after endpoints had been inspected, so this part is exploratory. Under them the full two-state classifica-




FIG. 2. (a) Mean waiting time of the perfect $4 \times L$ torus at $\lambda = 1.25$ against inverse coupling, sixteen runs per point; error bars are one standard error of the mean. Lines: Eq. (2) using move A only (solid) and moves A and B (dashed); nothing is fitted. (b) Share of vertices at each counting signature d at the first sampled $f \geq 0.5$, averaged over the thirty decays per size of the rerun in Table II. (c) Share of tori still waiting against time in units of the mean of its condition, pooled, on a logarithmic scale; the dashed line is the exponential e^{-x} of a memoryless wait. Blue: all 432 recorded first exits (twelve conditions of panel a, and the six cells of the direct count, T22). Orange: the selected 589 decay waits of the λ map with $f_{200} = 0$ and $W > 205$, shown as $W - 205$ (Methods). This selection does not establish that no earlier exit occurred. Not pre-registered.

tion holds at $N = 64, 96, 192$, while one $N = 144$ run fails because its state changes inside the final measuring window. The prespecified waiting and morphology criteria pass at all four sizes in both sets of thirty runs; a passing component test is not a passing combined classification.

A count is not a geometry. At $\lambda = 1.05$ all 120 saved endpoints of the map of Sec. VII have zero energy. Filling every square as a face and checking two faces per edge, a cyclic vertex link, connectedness, orientability and Euler characteristic zero certifies torus topology for all 120 graphs. This does not cover the other datasets or determine period lengths.

A separate equilibrium scan at $\lambda = 1.25$ has a single energy-distribution peak over the tested sizes and couplings and is INCONCLUSIVE under its stated rules. Neither that finite-size observation nor the kinetic conversion studied here establishes a thermodynamic first-

order transition.

VI. SEALED SYSTEM: THE RESERVOIR AND THE RELIC

Sealed, the released energy has to go somewhere, and the C stores of Sec. III are the only place (PL, Chs. 7 and 8).

Reading a temperature off the stores. At $\lambda = 1.25$, $\Delta H = -16\Delta S + 5\Delta X$ is integral. The initial T9 stores contain 0 or 12, and allowed exit costs include 12 and 25, with greatest common divisor one, so store energies have no common increment coarser than one; accessibility of all allocations is not thereby proved. If equilibrated with an approximately canonical graph, an unbounded store with levels $k\delta$ has a geometric energy distribution, whose mean u_D is tied to the coupling by

$$u_D(g) = \frac{\delta}{e^{\delta/g} - 1}, \quad g = \frac{\delta}{\ln(1 + \delta/u_D)}.$$

For $\delta = 1$, $u_D(3.5) = 3.024$, and $u_D = 0.5$ corresponds to $g = 0.910$. We report this inversion only as an equilibrium-conditional proxy: the CSV field `bath_T` is the mean store energy, not the inferred coupling, and store histograms were not saved, so their canonical distribution cannot be checked from these files.

Symmetric reverse moves support a uniform invariant measure on accessible states of graph plus stores at fixed total energy; they do not demonstrate equilibration. If all integer bath allocations at energy $M = E_{\text{tot}} - H(G) \geq 0$ are accessible, their number is $\Omega_C(M) = \binom{M+C-1}{C-1}$, and

TABLE II. The decay of the $4 \times L$ torus at $\lambda = 1.25$, $g = 1.5$ (first run / independent rerun, thirty decays each). “Largest piece” is the mean fraction of $d = 2$ vertices in their largest induced component. These are counting observables. Replica distributions and 95% bootstrap intervals are in the Supplemental Material; all 30 replicas per rerun size reach the half-conversion snapshot.

N	mean wait	CV	$d \in \{1, 2\}$ at 50% largest piece		
64	871 / 660	0.91 / 0.84	0.990 / 0.998	0.98 / 0.99	
96	913 / 928	1.10 / 1.02	0.990 / 0.988	0.95 / 0.96	
144	856 / 743	1.01 / 0.76	0.997 / 0.997	0.92 / 0.96	
192	1039 / 1428	0.87 / 1.19	0.995 / 0.995	0.85 / 0.92	

the graph marginal is weighted by this bath multiplicity. Changing C therefore changes the entropy as well as the thermal response.

Enough room, or not. The seed sits in one of C stores, and C sets the reservoir’s heat capacity ($C = 1, N/16, N/8, N/4, N/2, N, 2N$; $N = 64, 96, 192$; twenty runs each of 30,000 sweeps; 420 runs). With too little room the converted region melts into the random phase in 17 to 20 of 20 runs for $C \leq N/16$ and in 9 to 16 of 20 at $C = N/8$. At $C = N/4$ most runs end partly converted or partly melted. With enough room, $C \geq N/2$, the “sheet” counting criterion holds in 18 to 20 of 20; for perspective, 18/20 and 20/20 successes have exact 95% binomial intervals $[0.683, 0.988]$ and $[0.832, 1]$. The criterion requires $f \geq 0.9$, at least 90% at $d = 2$, and less than 5% at $d = 0$ or $d \geq 3$; it is not a topology test. The observed crossover lies between $N/4$ and $N/2$ on this coarse grid. A conditional equilibrium energy-balance estimate at crossover coupling g_m would be

$$C^* \simeq \frac{\varepsilon N + s - U_{\text{graph}}(g_m)}{u_D(g_m)}.$$

Taking $g_m \simeq 3.5$ and neglecting seed and residual graph energy gives $N/3.024$: an approximate energy budget, not a calibrated crossover prediction. Stores have no upper capacity, so a small bath heats; it does not fill up. No sustained stall with both signatures present was observed under the stated criterion and duration; this does not exclude equilibrium coexistence.

What remains: about one relic at every size. In the coldest box ($C = 2N$; T10, also 30,000 sweeps) a scrap of the old arrangement usually remains. A relic here is a connected component of the $d = 1$ induced subgraph in the final snapshot, classified further where wiring is available. It is the four-point loop above (14 units), its points lying along the torus and spanning one to four of the torus’s original columns, or occasionally a two-vertex twist (20 units). The number of relics per converted torus is 0.90 ± 0.31 , 0.85 ± 0.37 , 1.05 ± 0.22 and 0.95 ± 0.22 at $N = 64, 96, 192, 288$ (mean $\pm$ standard deviation over twenty runs each; a weighted fit against N gives a slope of 0.0003 ± 0.0003 per vertex, standard error). So the count is of order one, with no measurable growth, across the sizes we could reach. Where the relic sits shows no detected association with where the change began: the pre-registered test of whether it sits opposite the start or at it finds neither, and, checked afterwards, a randomisation test that rotates each relic through every column of its torus with the start held fixed places the observed mean distance at $p = 0.67$ and 0.59 ($N = 64$ and 96 ; 36 and 28 runs with one relic). These tests can miss other positional structure. Two cautions. The discrete coupling proxy ranges from 0.84 to 1.01 across the final baths, so even its conditional interpretation does not support a blanket “below $g = 1$ ” statement. And the relic count is a finite-duration observation, not a conclusion about the equilibrium density of defects.

VII. DEPENDENCE ON λ

T8 turns the knob: $\lambda = 1.05$ to 1.45 in steps of 0.05 at $g = 1.5$, using four sizes and thirty runs per cell (Table III); its $\lambda = 1.25$ comparison uses T7’s data. The registered criterion for a torus that waits (persistence) is that more than half the runs have $f_{200} < 0.25$. It holds from 1.05 to 1.35 at every size, and fails at 1.40 and 1.45 . This is an operational persistence boundary, not an energetic spinodal: move B still costs energy until $\lambda = 1.6$. Nor does Eq. (2) derive it. The equation gives a mean of 200 sweeps near $\lambda = 1.345$, but an exponential first exit with 50% survival at 200 would instead require a mean $200/\ln 2$, and f_{200} also includes growth.

In persistent cells, means at half conversion give 95–100% in the two counting signatures and 76–100% of flat-like vertices in the largest component. Release is fullest nearest $\lambda = 1$: the full-window-release fraction falls to 79% at 1.30 and 46% at 1.35 (the table retains this energy-window criterion rather than treating it as a topology test). The topology certificate for the $\lambda = 1.05$ endpoints is in Sec. V; PL, Ch. 9, walks through the map.

Three maps, each inconclusive. Three independently seeded maps fail their full prespecified acceptance criteria. In the thirty-run map (T8), the wait CV falls outside 0.7–1.3 in nine of 24 persistent cells, and the endpoint-energy criterion fails for two to five runs per size at $\lambda = 1.35$. In a 120-run map (T23), the sample-size-adjusted exponential CV criterion passes through 1.25 and in three sizes at 1.30, but the energy-window criterion fails at all sizes at 1.30 (thermal window averages need not match the energy of a saved endpoint). A second 120-run map (T24) tests endpoint validity and reports energy separately: all required saved graphs are valid, but the CV criterion fails at $(N, \lambda) = (64, 1.25), (64, 1.30), (192, 1.30), (144, 1.35)$, and validity is not evidence for full release. All three combined classifications are INCONCLUSIVE, while their counting-signature and dominant-component means pass through $\lambda = 1.30$. Full rules and cell results accompany the data.

Several early patches are common even when one component dominates later. From $\lambda = 1.30$ to 1.35 , the fraction of runs with several flat-like pieces at quarter conversion rises from 0.44 to 0.69 in T23 and 0.43 to 0.68 in T24, while the corresponding full-release or zero-energy fractions fall from 0.71 to 0.47 and 0.74 to 0.49, respectively (T24’s “FLAT” class means $H = 0$ only).

What the detector times. The decay detector first rests, taking the spread of S/N over a run’s first 200 sweeps as that run’s noise, and only then checks every five sweeps whether S/N has fallen more than three times that spread below its curled value (Table I; the “watch” of PL, Ch. 9). Its mean wait W and the first-exit prediction, compared over two decades in Fig. 3(a), are therefore distinct observables, for two reasons. First, nothing is recorded before sweep 205. For the idealized floor $W = \max(T, a)$ with exponential T of mean τ , $E[W] = a + \tau e^{-a/\tau}$, which at $a = 205$ gives 462 and 265 sweeps




FIG. 3. T8 at $g = 1.5$. (a) Detected mean waits in persistent cells; error bars are one standard error. Lines: Eq. (2) for first exits (solid) and the recording-floor approximation $205 + \tau_{AB}e^{-205/\tau_{AB}}$ (dashed); neither has fitted parameters. Shading marks failure of the operational persistence criterion. (b) Mean share in the two counting signatures at half conversion, and full-window-release fraction, pooled over sizes. Neither series certifies geometry.

at $\lambda = 1.30$ and 1.35 , against 403–505 and 280–350 measured (the dashed curve: a recording-floor approximation, not a calibrated detector prediction, and different from the conditional mean $E[T | T > a] = a + \tau$). Second, the threshold is the run’s own: a torus that begins to change inside the window widens its threshold and is caught late or never, which changes both the threshold and the sample selection (Methods).

The second effect is how we read the two exceptional T24 waits at $N = 64$: 20,270 sweeps ($\lambda = 1.25$) and 32,155 (1.30), or 23.98 and 76.74 times the predicted first-

exit means. They are not waits of a resting torus. The rows saved for those two runs record $f_{200} = 0.75$ and 0.875 and place the 25, 50 and 75% snapshots at sweep 205, the first check after the resting window: both tori had converted three quarters or more inside the window from which the detector estimates σ_{rest} . The threshold was correspondingly low, and W records when S/N at last fell below it, from endpoints that retained energy (window releases 0.744 of 1.000 and 0.563 of 1.200 per vertex).

TABLE III. T8 at $g = 1.5$. Ranges are cell means over $N = 64, 96, 144, 192$, thirty runs per cell; $\lambda = 1.25$ waits span both T7 sets, its release fraction uses the rerun. Signatures: share at $d \in \{1, 2\}$ at half conversion. Release: pooled fraction whose final window-averaged release is within 1% of ε , not a topology classification. At 1.05, each size has 30/30 and a 95% exact binomial interval $[0.884, 1]$; other intervals and denominators are in the Supplemental Material. A dash denotes failure of the persistence criterion.

λ	τ_{AB}	mean wait	signatures	release
1.05	8330	9617–11915	1.000	100%
1.10	4844	4154–6626	0.998–1.000	98%
1.15	2788	2191–3834	0.996–0.999	97%
1.20	1570	1132–1716	0.996–0.998	93%
1.25	845	660–1428	0.988–0.998	88%
1.30	419	403–505	0.980–0.984	79%
1.35	183	280–350	0.948–0.963	46%
1.40	68	—	0.911–0.924	31%
1.45	22	—	0.799–0.837	7%

Dedicated tail test (T38). A prespecified tail count asked whether long waits form a second population. With $w' = W - 200$ and $\hat{\tau} = \text{median}(w')/\ln 2$, it counts the waits k_{10} above $10\hat{\tau}$: TAIL for $k_{10} \geq 3$, NO TAIL for $k_{10} \leq 1$, otherwise UNCLEAR. At $\lambda = 1.25$ the result is UNCLEAR at $N = 64$ ($k_{10} = 2$; 3942 of 4000 runs recorded) and NO TAIL at $N = 192$ ($k_{10} = 1$; 1000 of 1000); at 1.30 it is TAIL at both sizes ($k_{10} = 14$, 3868 of 4000; and $k_{10} = 5$, 1000 of 1000). Two TAIL cells give the registered classification TWO POPULATIONS, which stands as scored; it does not establish two physical basins. A retrospective, exploratory reading of the saved rows, made after the data were seen (no new run; Methods and Supplemental Material), uses two recorded quantities that do not depend on the detector’s threshold, f_{200} and the sweep of the first sampled $f \geq 0.25$, and accounts for most of it by the resting window. (i) All 190 runs without a detector crossing have $f_{200} \geq 0.75$ and reach $f \geq 0.25$ at sweep 205: they converted inside the window and are the fastest runs, not long survivors. The tail statistic excludes them. (ii) The fastest detected waits are exits inside the window: of the runs detected within 20 sweeps of its end, 80–84% have $f_{200} > 0$, and the fraction of all runs with $f_{200} > 0$ is 0.144 and 0.146

at $\lambda = 1.25$ ($N = 64$ and 192) and 0.274 and 0.314 at 1.30 , against $1 - e^{-200/\tau_{AB}} = 0.211$ and 0.380 for a first exit before sweep 200 . (iii) The scale $\hat{\tau}$ is taken over all detected waits, those fast ones included: it is 750 , 671 , 267 and 238 sweeps in the four cells, against means of 917 , 858 , 443 and 446 for runs with $f_{200} = 0$. The cutoff $10\hat{\tau}$ is therefore only 5.7 – 8.9 τ_{AB} , and the reference expectation ne^{-10} does not apply to it. Measured against their own mean, the $f_{200} = 0$ runs have 0 , 0 , 3 and 0 waits beyond ten means, where 0.16 , 0.04 , 0.13 and 0.03 are expected.

That leaves three waits at $\lambda = 1.30$, $N = 64$ ($W = 4705$, 5000 and 5080). A registered replay of that cell from its seeds (T58) returns all $4,000$ saved rows unchanged. For the three runs it records every five-sweep block: before detection the square count never exceeds the torus's, no vertex reads $d = 0$, and the graph stays connected. One wait is the window again: an excursion inside it lowered the threshold, and three later exits that fell back went undetected. The other two are real by the counting signature: at every block the run read as the perfect torus, for $5,000$ and $4,670$ sweeps (11.9 and 11.1 τ_{AB}). Neither registered prediction for the three held (a hidden doubly curled state; a lowered threshold in all three). Those for the cell did. Timed with no detector, the first exits of the $4,000$ runs have mean 1.071 ± 0.017 τ_{AB} (four standard errors above the count; Methods), one lies beyond $10\tau_{AB}$ (0.18 expected), and no run ever exceeds the torus's square count. Described and not scored, two runs rest beyond $10\tau_{AB}$ after sweep 200 , where 0.13 are expected: rare, and unexplained.

VIII. DISCUSSION

The burp belongs to the $\lambda > 1$ family. At $\lambda = 1$, which is CQG, the flat, curled and 4-cube tori are degenerate and the available decrease $\varepsilon = 4(\lambda - 1)$ vanishes: the change has zero ideal-state energy difference. The same questions can still be asked in CQG itself, at $\lambda = 1$, of regions stuck in a different discrete arrangement inside a curved background [6], such as the allotrope of Ref. [3] (Fig. 9): is such a region metastable, and does it give way by a front with a definite release? The tools used here (exact barrier counting, front tracking, sealed baths) apply unchanged.

The strongest quantitative kinetic result is that counting moves predicts the mean first exit with nothing fitted, consistently with directly counted exits and with the omitted-channel correction stated. The energetic lower bound for sealed escape, 12 units at $\lambda = 1.25$, is certified over the tested sizes, while the available release grows with the system. Conversion snapshots support dominance by a connected flat-like region, but do not establish a unique seed or one moving boundary. Exact endpoint geometry is certified only for the stated T8 set. Finite-bath equilibration and period growth remain open. The detected-decay tails are largely accounted for, retrospec-

tively, by conversion inside the detector's resting window. A replay of one cell (Sec. VII) finds one more such case, two real waits that are rare and unexplained, and first exits within 7% of the count when timed without the detector. The inconclusive maps and the T38 classification are part of the result.

Named or interchangeable vertices. All runs here, like the published CQG simulations we know of [2, 3], treat vertices as named (labelled): two graphs that differ only by renaming vertices are different states. Giving each unlabelled class one Boltzmann weight is a different, legitimate modeling choice; quantum indistinguishability alone does not select this classical measure. A class has $(n!)^2/A(G)$ labelings, with $n = N/2$ and $A(G)$ its side-preserving automorphism count, so the corresponding target on labelled graphs is $A(G)e^{-H(G)/g}$, and the symmetric-proposal acceptance is

$$a(G \rightarrow G') = \min \left[1, \frac{A(G')}{A(G)} e^{-\Delta H/g} \right],$$

with the ratio inside the minimum (multiplying an already-clipped Metropolis probability is incorrect in general). Symmetric states then weigh more. Every energy and the certified energetic seed lower bound stay unchanged; the rates do not. The perfect curled torus has $2N$ symmetries and its neutral classes have N or $2N$, while the enumerated A-move targets in the $N = 64$, 96 calculation have two, so for those uphill moves the instantaneous acceptance is suppressed by $N/2$ to N . A rate suppression by cN corresponds to an effective free-energy increment $g \ln(cN)$, not a barrier proportional to N . Neutral-state occupancy, other exit channels and the quotient proposal also affect the lifetime; we supply no effective rate or decay simulation for this ensemble.

Not claimed: that this chain is the physical time evolution of anything; a first-order transition in the thermodynamic sense; a latent heat; isotropic decompactification; a unique moving front; a general large-size law beyond the enumerated and measured sizes; or an explanation for the cold-box defect count.

Outlook (motivation, not a result). This work is motivated by the hypothesis that space emerged from a specific ordered precursor through a sharp change of this kind, a lattice cousin of false-vacuum decay [8], releasing a definite energy and leaving relics behind. Nothing here tests that hypothesis beyond showing classical activated escape and energy-releasing conversion in a specified graph model. No tunneling or Lorentzian spacetime dynamics has been computed. Geometric periods, interfaces and attempt-resolved commitment histories would make the next tests more discriminating.

METHODS AND DATA

Sampler and validation. Simple bipartite switches preserve degree and the hard-core rule; invalid proposals are rejected. All clocks use sweeps of $2N$ attempts.

The stationary measure follows from symmetric proposals and Metropolis detailed balance. Exhaustive small-graph checks and canonical averages validate the sampler at $N = 16, 18$; ergodicity and equilibration at production sizes are not proved. The square-density comparison at $N = 160$ agrees with Ref. [2], Fig. 8a, to rms 0.005 but differs from Ref. [3], Fig. 3, by up to 0.41 between $g = 2$ and 6.3. Given the differing stated protocols discussed in the Introduction, this remains an unresolved reproduction issue, not a validation of the entire uncapped equilibrium curve.

Conversion protocol. The runner advances in five-sweep blocks, recording square counts and the threshold crossings in Table I. The main conversion stage stops at $f \geq 0.98$ or the configured cap; the settling stage follows even if that crossing was not reached. In T7b and T8, settling uses 600-sweep windows. The current stopping rule requires two consecutive window releases to agree within 0.005ϵ and the last to lie within 0.01ϵ of full release. Otherwise it runs to the configured settle cap. Defected endpoints can therefore exhaust the cap even when long-lived. T7 uses a fixed 300-sweep settling stretch. T7b's nominal main and settle caps are 30,000 sweeps each; T8 uses 100,000 for each. Recorded replay configurations extend selected T7 settling caps. The frozen configuration, not the nominal duration, identifies each replay. Saved final graphs are post-settling states, whereas the 25/50/75% histograms are taken during conversion.

Analysis design. Waiting and morphology criteria for the conversion maps, and the dedicated tail-count test, were fixed before their respective runs, as were the gate, labels and cell predictions of the T58 replay (one of its three target rows already implied a lowered threshold). Endpoint-energy criteria were partly chosen after inspecting data, so the energy classification is exploratory. The Arrhenius and seed scans, pooled distribution checks, topology certificates and interval estimates are also exploratory analyses, not independent confirmatory experiments. Experiment identifiers such as T8 and T38 connect the paper to the configurations and complete analysis record in the Supplemental Material. Our computational derivations remain unverified by an independent physicist.

Uncertainty and selection. Replicas, not vertices, are the independent units within a setting. We report 95% exact binomial intervals for success fractions, percentile bootstrap intervals for replica means (10,000 resamples, seed 20261009), and exponential-model intervals $[2 \sum T/\chi_{2m,0.975}^2, 2 \sum T/\chi_{2m,0.025}^2]$ for complete first-exit samples. The last are conditional on the continuous exponential approximation, not distribution-free bounds. For the T7b largest-component means, the intervals at $N = 64, 96, 144, 192$ are respectively $[0.983, 1]$, $[0.928, 0.986]$, $[0.906, 0.995]$, $[0.864, 0.971]$. The median is one in each cell, but 0, 1, 3 and 5 of the thirty individual runs have a fraction below 0.7 (0, 2, 4 and 9 in the first set). Thus a passing mean does not describe every replica. Map waiting summaries are restricted to

runs reaching $f \geq 0.75$; this selection can favor successful conversions. Counts of observed events, omissions and snapshot availability are given in the Supplemental Material.

Exploratory distribution checks. We divide each wait by its cell mean and compare the pooled sample with the unit exponential using the Kolmogorov–Smirnov distance D . A parametric bootstrap refits cell means in each of 10,000 simulated datasets. It reproduces the sampling grid and residual-time convention, not the full trajectory-dependent threshold. For the 432 first exits, two of eighteen individual cells have $p < 0.05$, both at $\lambda = 1.05$. The 240 attempt-counted T22 exits, normalized by Eq. (2) without fitted means, give $D = 0.039$, $p = 0.85$. The Arrhenius scan reuses a seed across couplings at each size and replica. The pooled bootstrap simulates cells independently, so its 432-wait p -value does not model this dependence and is nominal. T22 includes λ in its seed derivation. For detected waits, the unfiltered T7/T8 residual sample (865 waits with $W > 205$) gives $p = 0.03$; T7 alone gives $p < 10^{-3}$. Selecting T8 cells passing persistence, $f_{200} = 0$ and $W > 205$ gives 589 waits with $D = 0.030$, $p = 0.55$. Having $f_{200} = 0$ does not rule out an earlier exit and return. Pooling and fitting means can mask cell heterogeneity, and KS is not a decisive far-tail test. Three of 200 mixed first-exit and detected waits at $\lambda = 1.05$ exceed 75,000 sweeps; their post-hoc rarity calculation motivates an attempt-resolved audit.

Retrospective reading of the decay rows. The reading in Sec. VII uses only columns saved by the registered runs, `f_200` and the snapshot sweeps, through `scripts/read_wait_detector.py`. It was made after the T24 and T38 results were known, is exploratory, and changes no registered classification. A run is called perfect at sweep 200 when $f_{200} = 0$; this does not exclude an excursion that returned inside the window. The replay (T58) reruns registered configurations from their seeds with a read-only recorder; its gate requires every saved column to return unchanged. Its detector-free first exit is the first five-sweep block at which a run no longer has $S = 5N/4$, $X = N$ and every vertex at $d = 1$ (the basin of Table I, neutral classes included). An exit that returns within a block is not seen, which lengthens the measured wait by an amount not computed. No T38 waiting-graph files are in the frozen `results/` inventory, so in the three cells not replayed endpoints cannot reconstruct histories.

Companion documents. A plain-language edition (`paper_plain_language.pdf`, same repository) develops each concept with diagrams in the order of Secs. II–VII and links an interactive page evaluating Eqs. (1) and (2). Neither adds a result about this study; the edition's closing section previews later work.

Data and code. The source-data snapshot is available at RecreationalPhysics, commit 40b1ae68. The Supplemental Material provides a claim-to-script/configuration/data map, complete check outputs and content hashes in `release_manifest.json`, with the analysis and figure scripts. Configurations specify seeds

and durations; associated metadata record software versions but do not consistently identify the original execution commit. The retrospective-reading and replay scripts, the T58 configurations and results, and the plain-language edition postdate that snapshot and are in the same repository.

writing, and is responsible for the manuscript. We thank C. A. Trugenberger for correspondence.

ACKNOWLEDGMENTS

The author used Claude (Anthropic) and ChatGPT/Codex (OpenAI) to assist with code, analysis and

-
- [1] C. A. Trugenberger, J. High Energy Phys. **09** (2017) 045, arXiv:1610.05934.
 - [2] C. Kelly, C. A. Trugenberger, and F. Biancalana, Class. Quantum Grav. **36**, 125012 (2019), arXiv:1901.09870.
 - [3] C. A. Trugenberger, *Networks as the fundamental constituents of the universe*, J. Phys. Complex. **6**, 042001 (2025), arXiv:2512.17676.
 - [4] A. Gorsky and O. Valba, Phys. Rev. D **103**, 106013 (2021), arXiv:2101.04072.
 - [5] C. Kelly, *Emergent geometry and discrete curvature in random graphs*, Ph.D. thesis, Heriot-Watt University (2022), <https://www.ros.hw.ac.uk/handle/10399/4591>.
 - [6] C. A. Trugenberger, *Dark matter and dark energy in combinatorial quantum gravity*, arXiv:2409.09385.
 - [7] M. Creutz, Phys. Rev. Lett. **50**, 1411 (1983).
 - [8] S. Coleman, Phys. Rev. D **15**, 2929 (1977).